



Data Preprocessing & Visualization

Session : (Pandas)

The Dream...




... The Reality



“At the end of the day, some machine learning projects succeed and some fail. What makes the difference?”

Easily the most important factor is the features used.”

— Prof. Pedro Domingos



Exploratory Data Analysis (EDA)

Agenda:

1	Introduction to Pandas
2	Data Manipulation with Pandas
3	Practical Examples

Introduction to Pandas

Introduction to Pandas

What is Pandas?

Pandas is an open-source data analysis and data manipulation library built on top of Numpy. It provides high-level data structures and functions designed to make data analysis fast and easy in Python.

Pandas is widely used in data science, finance, statistics, and many other fields for tasks like data cleaning, transformation, and analysis.

Introduction to Pandas

Key Features of Pandas

- **Series and DataFrames:** Pandas introduces two primary data structures: Series (one-dimensional) and DataFrame (two-dimensional), which are used to store and manipulate data.
- **Data Alignment:** Pandas automatically aligns data by its labels, making it easy to perform arithmetic operations on data sets with different indexes.
- **Handling Missing Data:** Pandas provides methods for detecting, filling, and dropping missing data, which is essential for cleaning datasets.
- **Powerful Data Grouping and Aggregation:** Pandas supports grouping data, aggregating results, and applying custom functions using the `groupby()` functionality.
- **Flexible Data Reading/Writing:** Pandas can read from and write to various file formats, including CSV, Excel, SQL databases, and more.

Pandas

We will use pandas to:

- Read in data from Excel.
- Manipulate data in spreadsheet.
- Visualize data (we will also use another Python package called ggplot to do this).
- Filter and aggregate data from spreadsheet using SQL

Creating Pandas Data Structures

Creating a Pandas Series

```
import pandas as pd

# Creating a Series from a list
series = pd.Series([10, 20, 30, 40, 50])
print(series)
```

```
0    10
1    20
2    30
3    40
4    50
dtype: int64
```

Series

	apples
0	3
1	2
2	0
3	1

+

Series

	oranges
0	0
1	3
2	7
3	2

=

DataFrame

	apples	oranges
0	3	0
1	2	3
2	0	7
3	1	2

A Series is a one-dimensional labeled array capable of holding any data type. It is similar to a column in a spreadsheet.

Creating Pandas Data Structures

Creating a DataFrame from a Dictionary

```
import pandas as pd

[2] data = {'Name': ['Alice', 'Bob', 'Charlie'],
           'Age': [25, 30, 35],
           'City': ['New York', 'Los Angeles', 'Chicago']}

df = pd.DataFrame(data)
print(df)
```

```
⇒
```

	Name	Age	City
0	Alice	25	New York
1	Bob	30	Los Angeles
2	Charlie	35	Chicago

A DataFrame is a two-dimensional labeled data structure with columns of potentially different types. It is similar to a table in a database or a data frame in R.

	0	1	2	3	4		
Index Label	Name	Age	Marks	Grade	Hobby	Column Index	
0	S1	Joe	20	85.10	A	Swimming	
1	S2	Nat	21	77.80	B	Reading	
2	S3	Harry	19	91.54	A	Music	
3	S4	Sam	20	88.78	A	Painting	Row
4	S5	Monica	22	60.55	B	Dancing	

Annotations: Column, Row, Element/ Value/ Entry

Diagram illustrating a 2D array structure with row and column indices.

Column Label/ Header		0	1	2	3	4	Column Index
Index Label		Name	Age	Marks	Grade	Hobby	
0	S1	Joe	20	85.10	A	Swimming	Row
1	S2	Nat	21	77.80	B	Reading	
2	S3	Harry	19	91.54	A	Music	
3	S4	Sam	20	88.78	A	Painting	
4	S5	Monica	22	60.55	B	Dancing	
		Row Index		Column	Element/ Value/ Entry		

Missing Data

```
pd.read_csv?
```

`nrows : int, default None`

Number of rows of file to read. Useful for reading pieces of large files

`na_values` : scalar, str, list-like, or dict, default None

Additional strings to recognize as NA/NaN. If dict passed, specific per-column NA values. By default the following values are interpreted as NaN: `'', '#N/A', '#N/A N/A', '#NA', '-1.#IND', '-1.#QNAN', '-NaN', '-nan', '1.#IND', '1.#QNAN', 'N/A', 'NA', 'NULL', 'NaN', 'n/a', 'nan', 'null'`.

`keep_default_na : bool, default True`

If `na_values` are specified and `keep_default_na` is False the default NaN values are overridden, otherwise they're appended to.

Basic Features

```
import pandas as pd

df_grades = pd.read_csv("Grades_Short.csv")
df_grades.head(3)
```

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-

```
#dimension of df
df_grades.shape
```

(7, 9)

Think of this as a
list

```
#How each column is stored
df_grades.dtypes
```

```
Name          object
Previous_Part  float64
Participation1 int64
Mini_Exam1     float64
Mini_Exam2     int64
Participation2 int64
Mini_Exam3     float64
Final          float64
Grade          object
dtype: object
```

object = string

float64 = decimal

int64 = integer

Basic Features

column names



	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-

row names = index

```
#Get column names  
df_grades.columns
```


```
Index(['Name', 'Previous_Part', 'Participation1', 'Mini_Exam1', 'Mini_Exam2',  
      'Participation2', 'Mini_Exam3', 'Final', 'Grade'],  
      dtype='object')
```

```
#Get row names  
df_grades.index
```

```
RangeIndex(start=0, stop=7, step=1)
```


Viewing and Inspecting Dat

Viewing the First Few Rows



The image shows a Jupyter Notebook interface. At the top, there is a code cell with a play button icon and the text `df.head()`. Below the code cell, there is an output cell containing a table. The table has four columns: an index column, a 'Name' column, an 'Age' column, and a 'City' column. The first three rows of the table are visible, showing data for Alice, Bob, and Charlie.

	Name	Age	City
0	Alice	25	New York
1	Bob	30	Los Angeles
2	Charlie	35	Chicago

The `head()` method returns the first five rows of the DataFrame. This is useful for quickly examining the top of a dataset.

Viewing and Inspecting Data

Getting DataFrame Information

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3 entries, 0 to 2
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Name    3 non-null      object
1   Age     3 non-null      int64
2   City    3 non-null      object
dtypes: int64(1), object(2)
memory usage: 200.0+ bytes
```

The info() method provides a summary of the DataFrame, including the number of non-null entries and the data type of each column.

اخذت نظرة اولية على :

عدد الاعمدة - `df.head()`

عدد ال rows - `df.shape`

- `df.dtypes`

- `df.isnull().sum()`

	Name	Age	City
0	Alice	25	New York
1	Bob	30	Los Angeles
2	Charlie	35	Chicago

Viewing and Inspecting Dat

Describing Statistical Information

```
df.describe()
```



	Age
count	3.0
mean	30.0
std	5.0
min	25.0
25%	27.5
50%	30.0
75%	32.5
max	35.0

The describe() method generates descriptive statistics, such as mean, standard deviation, min, and max, for numerical columns.

```
df.head()
```



	Name	Age	City
0	Alice	25	New York
1	Bob	30	Los Angeles
2	Charlie	35	Chicago

"Whenever you open a new dataset, your **first 5 minutes** should include these five commands:

```
df.shape
```

```
df.columns
```

```
df.info()
```

```
df.describe()
```

```
df.head()
```

These will give you a **quick understanding** of the data's structure, size, types, and any missing values — before you do **any cleaning or analysis.**"

Selecting and Filtering Data

Selecting Columns

```
print(df['Name'])
```

```
0    Alice  
1     Bob  
2  Charlie  
Name: Name, dtype: object
```

You can select a single column from a DataFrame using the column name. This returns a Series object.

```
df.head()
```

	Name	Age	City
0	Alice	25	New York
1	Bob	30	Los Angeles
2	Charlie	35	Chicago

Selecting and Filtering Data

Filtering Rows Based on Condition

```
[7] filtered_df = df[df['Age'] > 30]
     print(filtered_df)
```

```
➡      Name  Age  City
2  Charlie  35  Chicago
```

You can filter rows based on a condition. Here, only rows where the 'Age' column is greater than 30 are selected.

df.head()



	Name	Age	City
0	Alice	25	New York
1	Bob	30	Los Angeles
2	Charlie	35	Chicago

Selecting and Filtering Data

Selecting Specific Rows and Columns

```
selected_data = df.loc[0:1, ['Name', 'City']]  
print(selected_data)
```

```
   Name      City  
0  Alice  New York  
1   Bob  Los Angeles
```

The `loc[]` method is used for label-based indexing to select specific rows and columns. Here, rows with indices 0 and 1 and columns 'Name' and 'City' are selected.

`loc[row , col]`

```
df.head()
```



	Name	Age	City
0	Alice	25	New York
1	Bob	30	Los Angeles
2	Charlie	35	Chicago

[go to code](#)

- Use `[]` for selecting columns
- Use `.loc[row_labels, column_labels]` for label-based indexing
- Use `.iloc[row_positions, column_positions]` for positional index

```
[37]: df
```

```
[37]:
```

	A	B	C
a	1	True	0.496714
b	2	True	-0.138264
c	3	False	0.647689

```
[21]: df.loc['a':'b', ['A', 'C']]
```

```
[21]:
```

	A	C
a	1	0.496714
b	2	-0.138264

```
[43]: df.iloc[0:2, [0,2]]
```

```
[43]:
```

	A	C
a	1	0.496714
b	2	-0.138264

Modifying Data

Adding a New Column

```
df['Salary'] = [70000, 80000, 90000]  
print(df)
```

	Name	Age	City	Salary
0	Alice	25	New York	70000
1	Bob	30	Los Angeles	80000
2	Charlie	35	Chicago	90000

A new column 'Salary' is added to the DataFrame with specified values. This is a common operation to add calculated fields or new data.

df.head()

	Name	Age	City
0	Alice	25	New York
1	Bob	30	Los Angeles
2	Charlie	35	Chicago

Modifying Data

Updating Column Values

```
df['Age'] = df['Age'] + 1  
print(df)
```

```
   Name  Age  City  Salary  
0  Alice   26 New York  70000  
1   Bob   31 Los Angeles  80000  
2 Charlie   36   Chicago  90000
```

This updates the values in the 'Age' column by adding 1 to each value. Pandas allows vectorized operations, making this efficient.

```
df.head()
```

	Name	Age	City
0	Alice	25	New York
1	Bob	30	Los Angeles
2	Charlie	35	Chicago

Modifying Data

Dropping Columns

```
df = df.drop('City', axis=1)  
print(df)
```

```
   Name  Age  Salary  
0  Alice   26   70000  
1   Bob   31   80000  
2 Charlie   36   90000
```

The `drop()` method removes the 'City' column from the DataFrame. The `axis=1` argument specifies that a column is being dropped.

axis{0 or 'index', 1 or 'columns'}, default 0

Whether to drop labels from the index (0 or 'index') or columns (1 or 'columns').

```
print(df)
```

```
   Name  Age  City  Salary  
0  Alice   26 New York  70000  
1   Bob   31 Los Angeles  80000  
2 Charlie   36   Chicago  90000
```

Try: axis 0

Handling Missing Data

Detecting Missing Values

```
df_with_nan = pd.DataFrame({'A': [1, 2, None], 'B': [4, None, 6]})  
print(df_with_nan.isnull())
```



	A	B
0	False	False
1	False	True
2	True	False

The `isnull()` method detects missing values in the DataFrame and returns a DataFrame of the same shape with True or False values.

```
[99]: df_with_nan= pd.DataFrame({'A':[1,2, None], 'B':[None, None, 6]})
df_with_nan
```

```
[99]:
```

	A	B
0	1.0	NaN
1	2.0	NaN
2	NaN	6.0

```
[101]: df_with_nan.isnull()
```

```
[101]:
```

	A	B
0	False	True
1	False	True
2	True	False

```
[109]: df_with_nan.isna()
```

```
[109]:
```

	A	B
0	False	True
1	False	True
2	True	False

```
[103]: df_with_nan.isnull().sum()
```

```
[103]: A    1
      B    2
      dtype: int64
```

```
•[105]: #total sum of null
df_with_nan.isnull().sum().sum()
```

```
[105]: 3
```

Handling Missing Data

Filling Missing Values

```
# Create DataFrame from dict
student_df = pd.DataFrame(student_dict , index=['a', 'b', 'c','d'] )
student_df
```

```
[99]:
```

	Name	Age	Marks
a	Joe	20.0	NaN
b	Ali	NaN	NaN
c	Baraa	40.0	95.0
d	Tala	35.0	80.0

```
[101]: student_df.fillna(0)
```

```
[101]:
```

	Name	Age	Marks
a	Joe	20.0	0.0
b	Ali	0.0	0.0
c	Baraa	40.0	95.0
d	Tala	35.0	80.0

The `fillna()` method replaces missing values with a specified value, such as 0. This is useful for cleaning datasets before analysis.

```
[133]: df_with_nan= pd.DataFrame({'A':[1,2, None], 'B':[None, None, 6]})
df_with_nan
```

```
[133]:
```

	A	B
0	1.0	NaN
1	2.0	NaN
2	NaN	6.0

```
•[137]: df_with_nan['A'] = df_with_nan['A'].fillna( df_with_nan['A'].mean())
```

```
[139]: df_with_nan
```

```
[139]:
```

	A	B
0	1.0	NaN
1	2.0	NaN
2	1.5	6.0

```
[131]: df_with_nan['B'] = df_with_nan['B'].fillna( df_with_nan['B'].mean() )
df_with_nan
```

```
[131]:
```

	A	B
0	1.0	6.0
1	2.0	6.0
2	1.5	6.0


```
# Python dict object
student_dict = {'Name': ['Joe', 'Ali', 'Baraa', "Tala"], 'Age': [20, None, 40, 35], 'Marks': [None, None, 95, 80]}

# Create DataFrame from dict
student_df = pd.DataFrame(student_dict , index=['a', 'b', 'c', 'd'] )
student_df
```

```
[99]:
```

	Name	Age	Marks
a	Joe	20.0	NaN
b	Ali	NaN	NaN
c	Baraa	40.0	95.0
d	Tala	35.0	80.0

```
[105]: student_df["Marks"].mean()
```

```
[105]: 87.5
```

```
[111]: student_df["Marks"].fillna( student_df["Marks"].mean() )
```

```
[111]:
```

a	87.5
b	87.5
c	95.0
d	80.0

Name: Marks, dtype: float64

Handling Missing Data

Dropping Rows with Missing Values

```
df_with_nan = pd.DataFrame({'A': [1, 2, None], 'B': [4, None, 6]})  
print(df_with_nan.isnull())
```

```
   A     B  
0  False False  
1  False  True  
2   True  False
```

```
cleaned_df = df_with_nan.dropna()  
print(cleaned_df)
```

```
   A     B  
0  1.0  4.0
```

The `dropna()` method removes rows with any missing values. This can be useful when working with datasets where missing values cannot be tolerated.

Try axis 1

Missing Data

A	B	C	D	E	F	G	H	I	J
Name	Previous_Par	Participation	Mini_Exam1	Mini_Exam2	Participation	Mini_Exam3	Final	Grade	Temp
Jake	32	1	19.5	20	1	10	33	A	-1
Joe	NA	1	20	16	1	14	32	A	23
Sol	31	1	22	13	1	13	34	A	34
Chris	30	-1	19	not available	1	12.5	33.5	A	72

```
df_missing = pd.read_csv("Data/Missing_Data.csv")
df_missing
```

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade	Temp
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A	-1
1	Joe	NaN	1	20.0	16	1	14.0	32.0	A	23
2	Sol	31.0	1	22.0	13	1	13.0	34.0	A	34
3	Chris	30.0	-1	19.0	not available	1	12.5	33.5	A	72

Not that different columns have different indicators for missing data.

Missing Data

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade	Temp
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A	-1
1	Joe	NaN	1	20.0	16	1	14.0	32.0	A	23
2	Sol	31.0	1	22.0	13	1	13.0	34.0	A	34
3	Chris	30.0	-1	19.0	not available	1	12.5	33.5	A	72

```
df_missing.dtypes
```

```
Name          object
Previous_Part  float64
Participation1 int64
Mini_Exam1     float64
Mini_Exam2     object
Participation2 int64
Mini_Exam3     float64
Final          float64
Grade          object
Temp          int64
dtype: object
```

na_values : Hashable, Iterable of Hashable or dict of {Hashable : Iterable},
optional

Additional strings to recognize as `NA`/`NaN`. If `dict` passed, specific per-column `NA` values. By default the following values are interpreted as `NaN`: " ", "#N/A", "#N/A N/A", "#NA", "-1.#IND", "-1.#QNAN", "-NaN", "-nan", "1.#IND", "1.#QNAN", "<NA>", "N/A", "NA", "NULL", "NaN", "None", "n/a", "nan", "null".

We can replace the missing data with a true NaN (right now everything is just a string).

Missing Data

```
df_missing = pd.read_csv("Data/Missing_Data.csv", \
                          na_values=["NaN", "not available"])
df_missing
```



List of strings specifying which values are missing.

Missing Data

```
df_missing = pd.read_csv("Data/Missing_Data.csv", \
                          na_values=["NaN", "not available"])
df_missing
```

List of strings specifying which values are missing.

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade	Temp
0	Jake	32.0	1	19.5	20.0	1	10.0	33.0	A	-1
1	Joe	NaN	1	20.0	16.0	1	14.0	32.0	A	23
2	Sol	31.0	1	22.0	13.0	1	13.0	34.0	A	34
3	Chris	30.0	-1	19.0	NaN	1	12.5	33.5	A	72

Special NaN value (from numpy package), which is not a string.

Missing Data

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade	Temp
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A	-1
1	Joe	NaN	1	20.0	16	1	14.0	32.0	A	23
2	Sol	31.0	1	22.0	13	1	13.0	34.0	A	34
3	Chris	30.0	-1	19.0	not available	1	12.5	33.5	A	72

We know “NaN” and “not available” are missing data points, but what about -1?

Missing Data

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade	Temp
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A	-1
1	Joe	NaN	1	20.0	16	1	14.0	32.0	A	23
2	Sol	31.0	1	22.0	13	1	13.0	34.0	A	34
3	Chris	30.0	-1	19.0	not available	1	12.5	33.5	A	72

We know “NaN” and “not available” are missing data points, but what about -1?

- For the Participation1 column the -1 is probably missing data.
- For the Temp column, the -1 is likely not missing data, since -1 is a valid temperature.

For each column, we can specify exactly which values correspond to missing data.

Missing Data

```
df_missing = pd.read_csv("Data/Missing_Data.csv", na_values={"Mini_Exam2" : "not available",\  
df_missing                                     "Participation1": -1})
```

Curly brackets

A diagram consisting of two black arrows pointing from the text 'Curly brackets' to the curly braces of the dictionary in the na_values parameter of the pandas read_csv function. The first arrow points to the opening curly brace '{', and the second arrow points to the closing curly brace '}'.

Missing Data

```
df_missing = pd.read_csv("Data/Missing_Data.csv", na_values={"Mini_Exam2" : "not available",\
                                                             "Participation1": -1})
df_missing
```

Column name as string



NaN value



Missing Data

```
df_missing = pd.read_csv("Data/Missing_Data.csv", na_values={"Mini_Exam2" : "not available",\
"Participation1": -1})
```

df_missing

Column name as string



NaN value



“For column Participation1, replace all -1s with a NaN.”

Missing Data

```
df_missing = pd.read_csv("Data/Missing_Data.csv", na_values={"Mini_Exam2" : "not available",\
                                                             "Participation1": -1})
```

```
df_missing
```

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade	Temp
0	Jake	32.0	1.0	19.5	20.0	1	10.0	33.0	A	-1
1	Joe	NaN	1.0	20.0	16.0	1	14.0	32.0	A	23
2	Sol	31.0	1.0	22.0	13.0	1	13.0	34.0	A	34
3	Chris	30.0	NaN	19.0	NaN	1	12.5	33.5	A	72

Notice that the -1 was replaced only in Participation1 column

select salary,salary*2 from emp

Grouping and Aggregation

Applying Custom Functions with apply()

```
def double_salary(x):  
    return x * 2  
  
df['Double Salary'] = df['Salary'].apply(double_salary)  
print(df)
```

	Department	Salary	Double Salary
0	HR	50000	100000
1	IT	60000	120000
2	HR	45000	90000
3	IT	80000	160000

The apply() method is used to apply a custom function to each element in a column. Here, it doubles the salary for each employee.

Grouping and Aggregation

```
[15] df = pd.DataFrame({'Department': ['HR', 'IT', 'HR', 'IT'],  
                      'Salary': [50000, 60000, 45000, 80000]})  
grouped = df.groupby('Department').mean()  
print(grouped)
```



	Salary
Department	
HR	47500.0
IT	70000.0

select avg(salary)
from users
group by Department

The groupby() method groups the data by a specified column and then calculates the mean salary for each department. This is commonly used for summarizing data.

Practical Examples

Calculating the Average Salary by Department

```
average_salary_by_dept = df.groupby('Department')['Salary'].mean()  
print("Average Salary by Department:\n")  
average_salary_by_dept.head()
```

➔ Average Salary by Department:

Salary	
Department	
Finance	87500.000000
HR	53000.000000
IT	70333.333333
Marketing	69000.000000

dtype: float64

```
print("First 5 Rows:\n")  
df.head()
```

➔ First 5 Rows:

	ID	Name	Department	Age	Salary	Joining_Date
0	1	John Doe	HR	28	50000	2020-05-21
1	2	Jane Smith	IT	34	75000	2019-06-12
2	3	Bob Johnson	Finance	45	85000	2018-07-15
3	4	Alice White	IT	29	72000	2021-01-04
4	5	Charlie Brown	Marketing	32	68000	2019-11-30

The `groupby()` method groups the data by the 'Department' column and calculates the mean salary for each department using `mean()`.

Selecting a Single Column

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A
5	Tarik	31.0	1	19.0	19	1	8.0	24.0	B
6	Malik	31.5	1	20.0	21	1	9.0	36.0	A

```
#Get Name column  
df_grades[ 'Name' ]
```

```
0    Jake  
1     Joe  
2   Susan  
3     Sol  
4   Chris  
5   Tarik  
6   Malik  
Name: Name, dtype: object
```

- Between square brackets, the column must be given as a string
- Outputs column as a series
 - A series is a one dimensional dataframe..more on this in the slicing section

Selecting a Single Column

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A
5	Tarik	31.0	1	19.0	19	1	8.0	24.0	B
6	Malik	31.5	1	20.0	21	1	9.0	36.0	A

```
#Get Name column  
df_grades.Name
```

```
0    Jake  
1     Joe  
2   Susan  
3     Sol  
4   Chris  
5   Tarik  
6   Malik  
Name: Name, dtype: object
```

- Exactly equivalent way to get Name column
 - + : don't have to type brackets or quotes
 - -: won't generalize to selecting multiple columns,, won't work if column names have spaces, can't create new columns this way

Selecting Multiple Columns

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A
5	Tarik	31.0	1	19.0	19	1	8.0	24.0	B
6	Malik	31.5	1	20.0	21	1	9.0	36.0	A

```
#Select multiple columns  
df_grades[["Name", "Grade"]]
```

	Name	Grade
0	Jake	A
1	Joe	A
2	Susan	A-
3	Sol	A
4	Chris	A
5	Tarik	B
6	Malik	A

- List of strings, which correspond to column names.
- You can select as many column as you want.
- Column don't have to be contiguous.

Storing Result

```
#Print the column  
df_grades[ "Name" ]
```

```
0    Jake  
1     Joe  
2   Susan  
3     Sol  
4   Chris  
5   Tarik  
6   Malik  
Name: Name, dtype: object
```

```
#Store the column  
names= df_grades[ "Name" ]  
names
```

```
0    Jake  
1     Joe  
2   Susan  
3     Sol  
4   Chris  
5   Tarik  
6   Malik  
Name: Name, dtype: object
```

The variable name stores a series

Why store a slice?

- We might want/have to do our analysis in steps.
 - Less error prone
 - More readable

Slicing a Series

Slice/index through the index, which is usually numbers



```
names= df_grades[ "Name" ]  
names
```

```
0    Jake  
1     Joe  
2   Susan  
3     Sol  
4   Chris  
5   Tarik  
6   Malik  
Name: Name, dtype: object
```

Picking out single element

```
names[0]
```

```
'Jake'
```

Slicing a Series

Slice/index through the index, which is usually numbers

```
names= df_grades[ "Name" ]
names
```



0	Jake
1	Joe
2	Susan
3	Sol
4	Chris
5	Tarik
6	Malik

Name: Name, dtype: object

Picking out single element

```
names[0]
```

'Jake'

Contiguous slice

non_inclusive



```
names[1:4]
```

1	Joe
2	Susan
3	Sol

Name: Name, dtype: object

Slicing a Series

Slice/index through the index, which is usually numbers

```
names= df_grades[ "Name" ]
names
```



0	Jake
1	Joe
2	Susan
3	Sol
4	Chris
5	Tarik
6	Malik

Name: Name, dtype: object

Picking out single element

```
names[0]
```

'Jake'

Contiguous slice

```
names[1:4]
```

1	Joe
2	Susan
3	Sol

Name: Name, dtype: object

Arbitrary slice

```
names[[1,2,4]]
```

1	Joe
2	Susan
4	Chris

Name: Name, dtype: object

Slicing a Data Frame

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A
5	Tarik	31.0	1	19.0	19	1	8.0	24.0	B
6	Malik	31.5	1	20.0	21	1	9.0	36.0	A

- There are a few ways to pick slice a data frame, we will use the .loc method.
- Access elements through the index labels column names
 - We will see how to change both of these labels later on

Slicing a Data Frame

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A
5	Tarik	31.0	1	19.0	19	1	8.0	24.0	B
6	Malik	31.5	1	20.0	21	1	9.0	36.0	A

- Pick a single value out.

Column name (string)

Index label
(number)

```
first_name = df_grades.loc[0, "Name"]  
first_name  
'Jake'
```


Slicing a Data Frame

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A
5	Tarik	31.0	1	19.0	19	1	8.0	24.0	B
6	Malik	31.5	1	20.0	21	1	9.0	36.0	A

- Pick out entire row:

```
first_row = df_grades.loc[0,:]
first_row
```

“pick out all
columns”

```
Name      Jake
Previous_Part      32
Participation1      1
Mini_Exam1      19.5
Mini_Exam2      20
Participation2      1
Mini_Exam3      10
Final      33
Grade      A
Name: 0, dtype: object
```

first_row is a series

Slicing a Data Frame

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A
5	Tarik	31.0	1	19.0	19	1	8.0	24.0	B
6	Malik	31.5	1	20.0	21	1	9.0	36.0	A

- Pick out contiguous chunk:

Endpoints are inclusive!

```
slice_one = df_grades.loc[0:2, "Name": "Mini_Exam2"]  
slice_one
```

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2
0	Jake	32.0	1	19.5	20
1	Joe	32.0	1	20.0	16
2	Susan	30.0	1	19.0	19

df_grades.iloc[0:3, 0:5]

Slicing a Data Frame

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A
5	Tarik	31.0	1	19.0	19	1	8.0	24.0	B
6	Malik	31.5	1	20.0	21	1	9.0	36.0	A

- Pick out arbitrary chunk:

```
slice_two = df_grades.loc[[0,2,3], ["Name", "Grade"]]  
slice_two
```

	Name	Grade
0	Jake	A
2	Susan	A-
3	Sol	A

Built in Functions

```
import pandas as pd

df_grades = pd.read_csv("Data/Grades_Short.csv")
df_grades
```

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A
5	Tarik	31.0	1	19.0	19	1	8.0	24.0	B
6	Malik	31.5	1	20.0	21	1	9.0	36.0	A

How do I compute the average score on the final?

Built in Functions

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A
5	Tarik	31.0	1	19.0	19	1	8.0	24.0	B
6	Malik	31.5	1	20.0	21	1	9.0	36.0	A

How do I compute the average score on the final?

```
#Print out  
df_grades.Final.mean()
```

```
32.214285714285715
```

```
#Store  
avg_final = df_grades.Final.mean()  
avg_final
```

```
32.214285714285715
```

Built in mean() method

Built in Functions

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A
5	Tarik	31.0	1	19.0	19	1	8.0	24.0	B
6	Malik	31.5	1	20.0	21	1	9.0	36.0	A

How do I compute the highest Mini Exam 1 score?

```
max_mini_1 = df_grades["Mini_Exam1"].max()  
max_mini_1
```

22.0

Built in Functions

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A
5	Tarik	31.0	1	19.0	19	1	8.0	24.0	B
6	Malik	31.5	1	20.0	21	1	9.0	36.0	A

I can actually get all key stats for *numeric* columns at once with the describe() method:

```
summary_df = df_grades.describe()  
summary_df
```

summary_df is a
dataframe!

	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final
count	7.000000	7.0	7.000000	7.000000	7.0	7.000000	7.000000
mean	31.071429	1.0	19.785714	17.857143	1.0	11.000000	32.214286
std	0.838082	0.0	1.074598	2.734262	0.0	2.217356	3.828154
min	30.000000	1.0	19.000000	13.000000	1.0	8.000000	24.000000
25%	30.500000	1.0	19.000000	16.500000	1.0	9.500000	32.500000
50%	31.000000	1.0	19.500000	19.000000	1.0	10.500000	33.000000
75%	31.750000	1.0	20.000000	19.500000	1.0	12.750000	33.750000
max	32.000000	1.0	22.000000	21.000000	1.0	14.000000	36.000000

Built in Functions

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A
5	Tarik	31.0	1	19.0	19	1	8.0	24.0	B
6	Malik	31.5	1	20.0	21	1	9.0	36.0	A

I can actually get all key stats for *numeric* columns at once with the describe() method:

```
summary_df = df_grades.describe()  
summary_df[["Final", "Mini_Exam3"]]
```

	Final	Mini_Exam3
count	7.000000	7.000000
mean	32.214286	11.000000
std	3.828154	2.217356
min	24.000000	8.000000
25%	32.500000	9.500000
50%	33.000000	10.500000
75%	33.750000	12.750000
max	36.000000	14.000000

Built in Functions

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A
5	Tarik	31.0	1	19.0	19	1	8.0	24.0	B
6	Malik	31.5	1	20.0	21	1	9.0	36.0	A

Other useful built in methods:

```
df_grades["Grade"].value_counts()
```

```
A      5
A-     1
B       1
Name: Grade, dtype: int64
```

value_count(): Gives a count of the number of times each unique value appears in the column. Returns a series where indices are the unique column values.

Built in Functions

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A
5	Tarik	31.0	1	19.0	19	1	8.0	24.0	B
6	Malik	31.5	1	20.0	21	1	9.0	36.0	A

Other useful built in methods:

```
counts = df_grades["Grade"].value_counts()  
counts
```

```
A      5  
A-     1  
B      1  
Name: Grade, dtype: int64
```

```
counts["A"]
```

5

value_count(): Gives a count of the number of times each unique value appears in the column. Returns a series where indices are the unique column values.

Built in Functions

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A
5	Tarik	31.0	1	19.0	19	1	8.0	24.0	B
6	Malik	31.5	1	20.0	21	1	9.0	36.0	A

Other useful built in methods:

```
df_grades["Grade"].unique()
```

```
array(['A', 'A-', 'B'], dtype=object)
```

```
unique_values = df_grades["Grade"].unique()  
unique_values[0]
```

```
'A'
```

```
len(unique_values)
```

```
3
```

unique(): Returns an array of all of the unique values.

Attributes vs. Methods

When do I put a ()?

```
#Get dimensions  
df_grades.shape
```

```
(7, 9)
```

```
#Get column types  
df_grades.dtypes
```

```
Name          object  
Previous_Part  float64  
Participation1 int64  
Mini_Exam1     float64  
Mini_Exam2     int64  
Participation2 int64  
Mini_Exam3     float64  
Final          float64  
Grade          object  
dtype: object
```

```
#Get first 5 rows  
df_grades.head()
```

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A

```
#Compute mean of finals column  
df_grades["Final"].mean()
```

```
32.214285714285715
```


Attributes vs. Methods

When do I put a ()?

dataframe attributes

```
#Get dimensions  
df_grades.shape
```

```
(7, 9)
```

```
#Get column types  
df_grades.dtypes
```

```
Name          object  
Previous_Part  float64  
Participation1 int64  
Mini_Exam1     float64  
Mini_Exam2     int64  
Participation2 int64  
Mini_Exam3     float64  
Final          float64  
Grade          object  
dtype: object
```

dataframe methods

```
#Get first 5 rows  
df_grades.head()
```

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A

```
#Compute mean of finals column  
df_grades["Final"].mean()
```

```
32.214285714285715
```

Attributes vs. Methods

When do I put a ()?

dataframe attributes

```
#Get dimensions  
df_grades.shape
```

```
(7, 9)
```

```
#Get column types  
df_grades.dtypes
```

```
Name          object  
Previous_Part  float64  
Participation1 int64  
Mini_Exam1     float64  
Mini_Exam2     int64  
Participation2 int64  
Mini_Exam3     float64  
Final          float64  
Grade          object  
dtype: object
```

Features of dataframe

dataframe methods

```
#Get first 5 rows  
df_grades.head()
```

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A

```
#Compute mean of finals column  
df_grades["Final"].mean()
```

```
32.214285714285715
```

Require computation for output

Creating New Columns

```
import pandas as pd

df_grades = pd.read_csv("Data/Grades_Short.csv")
df_grades.head()
```

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A

Let's create a useless new column of all 1s:

```
df_grades["new_column"]=1
df_grades.head()
```

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade	new_column
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A	1
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A	1
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-	1
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A	1
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A	1

Creating New Columns

```
import pandas as pd
```

```
df_grades = pd.read_csv("Data/Grades_Short.csv")  
df_grades.head()
```

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A

33/36

32/36

We can also create column as function of other column. The Final was worth 36 points, let's create a column for each student's percentage.

```
df_grades["Final_Percentage"] = df_grades["Final"] / 36  
df_grades.head()
```

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade	new_column	Final_Percentage
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A	1	0.916667
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A	1	0.888889
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-	1	0.916667
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A	1	0.944444
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A	1	0.930556

Deleting Columns

```
df_grades.head()
```

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade	new_column	Final_Percentage	Part_perc
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A	1	0.916667	1.0
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A	1	0.888889	1.0
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-	1	0.916667	1.0
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A	1	0.944444	1.0
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A	1	0.930556	1.0

```
#Delete single column  
del df_grades["Final_Percentage"]  
df_grades.head()
```

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade	new_column	Part_perc
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A	1	1.0
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A	1	1.0
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-	1	1.0
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A	1	1.0
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A	1	1.0

Deleting Columns

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade	new_column	Part_perc
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A	1	1.0
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A	1	1.0
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-	1	1.0
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A	1	1.0
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A	1	1.0

```
#Delete multiple columns
```

```
df_grades.drop(["new_column", "Part_perc"], axis=1, inplace = True)
```

```
df_grades.head()
```

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A
1	Joe	32.0	1	20.0	16	1	14.0	32.0	A
2	Susan	30.0	1	19.0	19	1	10.5	33.0	A-
3	Sol	31.0	1	22.0	13	1	13.0	34.0	A
4	Chris	30.0	1	19.0	17	1	12.5	33.5	A

Comparing Approaches

Approach 1:

- Does a global search and replace in all columns.

```
df_missing = pd.read_csv("Data/Missing_Data.csv", \
                          na_values=[ "NaN", "not available" ])
df_missing
```

Approach 2:

- Allows you to specify column by column the values that should be replaced with NaN.

[illegible]

Benefiting of Having NaNs

- Have common symbol for where there is missing data
 - Good for you and good for others looking at your code/data
 - These entries will be ignored if you try to compute means of columns with NaNs.
- We can easily get rid of column/rows with missing data
- We can easily replace the missing values with the mean of the column, for example.

IsNull() Method

- The isnull() method lets you check where the NaNs are:

```
df = pd.read_csv("Data/Missing_Data.csv", na_values=["NaN", -1, "not available"])
df
```

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade	Temp
0	Jake	32.0	1.0	19.5	20.0	1	10.0	33.0	A	NaN
1	Joe	NaN	1.0	20.0	16.0	1	14.0	32.0	A	23.0
2	Sol	31.0	1.0	22.0	13.0	1	13.0	34.0	A	34.0
3	Chris	30.0	NaN	19.0	NaN	1	12.5	33.5	A	72.0

```
#Using isnull()
df.isnull()
```

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade	Temp
0	False	False	False	False	False	False	False	False	False	True
1	False	True	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False
3	False	False	True	False	True	False	False	False	False	False


```
[166]: df_WithNotNan=pd.read_csv("Data/Missing_Data.csv" )
df_WithNotNan
```

```
[166]:
```

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade	Temp
0	Jake	32.0	1	19.5	20	1	10.0	33.0	A	-1
1	Joe	NaN	1	20.0	16	1	14.0	32.0	A	23
2	Sol	31.0	1	22.0	13	1	13.0	34.0	A	34
3	Chris	30.0	-1	19.0	not available	1	12.5	33.5	A	72

```
[168]: df_WithNotNan.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4 entries, 0 to 3
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Name             4 non-null     object
1   Previous_Part    3 non-null     float64
2   Participation1    4 non-null     int64
3   Mini_Exam1       4 non-null     float64
4   Mini_Exam2       4 non-null     object
5   Participation2    4 non-null     int64
6   Mini_Exam3       4 non-null     float64
7   Final            4 non-null     float64
8   Grade            4 non-null     object
9   Temp             4 non-null     int64
dtypes: float64(4), int64(3), object(3)
memory usage: 452.0+ bytes
```

```
[170]: df_WithNotNan.isnull().sum()
```

```
[170]: Name             0
Previous_Part      1
Participation1     0
Mini_Exam1         0
Mini_Exam2         0
Participation2     0
Mini_Exam3         0
Final              0
Grade              0
Temp              0
dtype: int64
```

```
[172]: df_WithNotNan["Mini_Exam2"] = pd.to_numeric(df_WithNotNan["Mini_Exam2"],errors = "coerce")
df_WithNotNan
```

```
[172]:
```

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade	Temp
0	Jake	32.0	1	19.5	20.0	1	10.0	33.0	A	-1
1	Joe	NaN	1	20.0	16.0	1	14.0	32.0	A	23
2	Sol	31.0	1	22.0	13.0	1	13.0	34.0	A	34
3	Chris	30.0	-1	19.0	NaN	1	12.5	33.5	A	72

```
[176]: df_WithNotNan.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4 entries, 0 to 3
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Name            4 non-null     object
1   Previous_Part   3 non-null     float64
2   Participation1  4 non-null     int64
3   Mini_Exam1     4 non-null     float64
4   Mini_Exam2     3 non-null     float64
5   Participation2  4 non-null     int64
6   Mini_Exam3     4 non-null     float64
7   Final          4 non-null     float64
8   Grade          4 non-null     object
9   Temp           4 non-null     int64
```

IsNull() Method

- The isnull() method lets you check where the NaNs are:

```
▼ #Using isnull()  
df.isnull()
```

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade	Temp
0	False	False	False	False	False	False	False	False	False	True
1	False	True	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False
3	False	False	True	False	True	False	False	False	False	False

```
▼ #Remember Booleans are just 0s and 1s.  
#Check how many NaNs are in each column  
df.isnull().sum()
```

```
Name          0  
Previous_Part  1  
Participation1 1  
Mini_Exam1    0  
Mini_Exam2    1  
Participation2 0  
Mini_Exam3    0  
Final         0  
Grade         0  
Temp         1  
dtype: int64
```

Dropna() Method

```
df_missing = pd.read_csv("Data/Missing_Data.csv", na_values=["NaN", "not available", \
-1])
df_missing
```

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade	Temp
0	Jake	32.0	1.0	19.5	20.0	1	10.0	33.0	A	NaN
1	Joe	NaN	1.0	20.0	16.0	1	14.0	32.0	A	23.0
2	Sol	31.0	1.0	22.0	13.0	1	13.0	34.0	A	34.0
3	Chris	30.0	NaN	19.0	NaN	1	12.5	33.5	A	72.0

How do I get rid of all rows with NaN?

Practical Examples

Practical Examples

We'll load this CSV file into a Pandas DataFrame and perform various data manipulation tasks



```
import pandas as pd

# Load the CSV file into a Pandas DataFrame
df = pd.read_csv('employees.csv')
```

Practical Examples

Viewing the First Few Rows of the DataFrame

```
print("First 5 Rows:\n")  
df.head()
```

First 5 Rows:

	ID	Name	Department	Age	Salary	Joining_Date
0	1	John Doe	HR	28	50000	2020-05-21
1	2	Jane Smith	IT	34	75000	2019-06-12
2	3	Bob Johnson	Finance	45	85000	2018-07-15
3	4	Alice White	IT	29	72000	2021-01-04
4	5	Charlie Brown	Marketing	32	68000	2019-11-30

The `head()` method displays the first five rows of the DataFrame, which is useful for getting a quick look at the data.

Practical Examples

Displaying Basic Information About the DataFrame

```
▶ print("DataFrame Information:\n")  
df.info()
```

↔ DataFrame Information:

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 10 entries, 0 to 9  
Data columns (total 6 columns):  
#   Column          Non-Null Count  Dtype  ---  
0   ID              10 non-null    int64  ---  
1   Name            10 non-null    object ---  
2   Department      10 non-null    object ---  
3   Age             10 non-null    int64  ---  
4   Salary          10 non-null    int64  ---  
5   Joining_Date    10 non-null    object ---  
dtypes: int64(3), object(3)  
memory usage: 608.0+ bytes
```

The `info()` method provides a summary of the DataFrame, including the number of non-null values, data types, and memory usage.

Practical Examples

Calculating Basic Statistics for Numerical Columns

```
print("Statistical Summary:\n")  
df.describe()
```

Statistical Summary:

	ID	Age	Salary
count	10.00000	10.000000	10.000000
mean	5.50000	35.300000	68300.000000
std	3.02765	7.958922	13106.826042
min	1.00000	26.000000	50000.000000
25%	3.25000	29.250000	57250.000000
50%	5.50000	33.000000	69000.000000
75%	7.75000	40.250000	74250.000000
max	10.00000	50.000000	90000.000000

The `describe()` method generates descriptive statistics such as mean, standard deviation, and percentiles for numerical columns.

```
import pandas as pd
```

```
# We pass a dict of {column name: column values}
```

```
df = pd.DataFrame({'X': [78, 85, 96, 80, 86], 'Y': [84, None, None, None, None], 'Z': [86, 97, 96, 72, None]});  
df
```

```
[181]:
```

	X	Y	Z
0	78	84.0	86.0
1	85	NaN	97.0
2	96	NaN	96.0
3	80	NaN	72.0
4	86	NaN	NaN

```
[183]: df.describe()
```

```
[183]:
```

	X	Y	Z
count	5.0	1.0	4.000000
mean	85.0	84.0	87.750000
std	7.0	NaN	11.615363
min	78.0	84.0	72.000000
25%	80.0	84.0	82.500000
50%	85.0	84.0	91.000000
75%	86.0	84.0	96.250000

Practical Examples

Filtering Employees by IT Department

```
▶ it_employees = df[df['Department'] == 'IT']  
print("IT Department Employees:\n")  
it_employees.head()
```

➞ IT Department Employees:

	ID	Name	Department	Age	Salary	Joining_Date	
1	2	Jane Smith	IT	34	75000	2019-06-12	
3	4	Alice White	IT	29	72000	2021-01-04	
7	8	Linda Taylor	IT	26	64000	2022-02-10	

This code filters the DataFrame to show only employees who work in the IT department using a condition on the 'Department' column.

Practical Examples

Adding a New Column for Salary After a 5% Raise

```
df['Salary_After_Raise'] = df['Salary'] * 1.05  
print("DataFrame with Salary After Raise:\n" )  
df.head()
```

→ DataFrame with Salary After Raise:

	ID	Name	Department	Age	Salary	Joining_Date	Salary_After_Raise
0	1	John Doe	HR	28	50000	2020-05-21	52500.0
1	2	Jane Smith	IT	34	75000	2019-06-12	78750.0
2	3	Bob Johnson	Finance	45	85000	2018-07-15	89250.0
3	4	Alice White	IT	29	72000	2021-01-04	75600.0
4	5	Charlie Brown	Marketing	32	68000	2019-11-30	71400.0

A new column is added to the DataFrame by multiplying the 'Salary' column by 1.05 to reflect a 5% salary increase.

Practical Examples

Sorting Employees by Age in Descending Order

```
sorted_by_age = df.sort_values(by='Age', ascending=False)
print("Employees Sorted by Age (Descending):\n")
sorted_by_age.head()
```

Employees Sorted by Age (Descending):

	ID	Name	Department	Age	Salary	Joining_Date	Salary_After_Raise
6	7	Mike Davis	Finance	50	90000	2016-08-19	94500.0
2	3	Bob Johnson	Finance	45	85000	2018-07-15	89250.0
5	6	Eva Green	HR	41	54000	2017-04-22	56700.0
8	9	James Wilson	Marketing	38	70000	2020-09-08	73500.0
1	2	Jane Smith	IT	34	75000	2019-06-12	78750.0

https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.sort_values.html

The `sort_values()` method sorts the DataFrame by the 'Age' column in descending order, showing the oldest employees first.

Practical Examples

Calculating the Number of Employees in Each Department

```
▶ employee_count_by_dept = df['Department'].value_counts()  
print("Number of Employees by Department:\n")  
employee_count_by_dept.head()
```

⇒ Number of Employees by Department:

count	
Department	
HR	3
IT	3
Finance	2
Marketing	2

dtype: int64

The `value_counts()` method counts the number of occurrences of each unique value in the 'Department' column, showing how many employees are in each department.

Practical Examples

Finding Employees Who Joined After 2020

```
df['Joining_Date'] = pd.to_datetime(df['Joining_Date'])  
recent_joins = df[df['Joining_Date'] > '2020-01-01']  
print("Employees Joined After 2020:\n")  
recent_joins.head()
```

Employees Joined After 2020:

	ID	Name	Department	Age	Salary	Joining_Date	Salary_After_Raise
0	1	John Doe	HR	28	50000	2020-05-21	52500.0
3	4	Alice White	IT	29	72000	2021-01-04	75600.0
7	8	Linda Taylor	IT	26	64000	2022-02-10	67200.0
8	9	James Wilson	Marketing	38	70000	2020-09-08	73500.0
9	10	Sarah Adams	HR	30	55000	2021-03-18	57750.0

This converts the 'Joining_Date' column to datetime format and filters the DataFrame to find employees who joined after January 1, 2020.

Practical Examples


Saving the Filtered DataFrame to a New CSV File

```
recent_joins.to_csv('recent_joins.csv', index=False)  
print("Filtered data saved to recent_joins.csv.")
```

Filtered data saved to recent_joins.csv.

The `to_csv()` method saves the filtered DataFrame (`recent_joins`) to a new CSV file named `'recent_joins.csv'`. The `index=False` argument excludes the row index from the CSV file.

Summary



Using Apply

Let's see how we can use apply with a custom function on a single column:

```
def Get_Gender_Bool(gender):  
    if gender == "male":  
        return 1  
    else:  
        return 0
```

```
df_age_sex["Bool_Male"] = df_age_sex.Sex.apply(Get_Gender_Bool)  
df_age_sex
```

	Age	Sex	Bool_Male
0	34.5	male	1
1	47.0	female	0
2	62.0	male	1
3	27.0	male	1
4	22.0	female	0



Calls Get_Gender_Bool once for each row.

```
•[304]: def getResult(x):  
        if x == "A":  
            return "Excellent"  
        else:  
            return "Good"
```

```
[298]: df["Result"] = df["Grade"].apply( lambda x : "Excellent" if x == "A" else "Good" )
```

```
•[306]: df["Result1"] = df["Grade"].apply( getResult )
```



```
[308]: df
```

```
[308]:
```

	Name	Previous_Part	Participation1	Mini_Exam1	Mini_Exam2	Participation2	Mini_Exam3	Final	Grade	Grade Num	Result	Result1
0	Jake	32.0	1	19.5	20.0	1	10.0	33.0	A	A	Excellent	Excellent
1	?	32.0	1	20.0	16.0	1	14.0	32.0	A	A	Excellent	Excellent
2	Susan	30.0	1	19.0	19.0	1	10.5	33.0	A-	A-	Good	Good
3	Sol	31.0	1	22.0	13.0	1	13.0	34.0	A	A	Excellent	Excellent
4	Chris	30.0	1	19.0	17.0	1	12.5	33.5	A	A	Excellent	Excellent
5	Tarik	31.0	1	19.0	19.0	1	8.0	24.0	B	B	Good	Good

Using Apply

Last one:

```
def Get_Title(name):  
  
    parsed_name = name.split(" ")  
    title = parsed_name[1]  
  
    return title
```

```
df_name_age["Title"] = df_name_age.Name.apply(Get_Title)  
df_name_age
```

	Name	Age	Title
0	Kelly, Mr. James	34.5	Mr.
1	Wilkes, Mrs. James (Ellen Needs)	47.0	Mrs.
2	Myles, Mr. Thomas Francis	62.0	Mr.
3	Wirz, Mr. Albert	27.0	Mr.
4	Hirvonen, Mrs. Alexander (Helga E Lindqvist)	22.0	Mrs.

Data Preprocessing

Feature Transformations

- Data Cleaning or Cleansing
- Work with Missing data
- Work with Categorical data
- Detect and Handle Outliers
- Split data to Train and Test Sets
- Deal with Imbalanced classes
- Feature Scaling

Feature Engineering and Extraction

- Domain knowledge features (age , wh, **BMI**)
- Date and Time features
- String operations
- Web Data
- Geospatial features

Feature Transformations

1. Data Cleaning or Cleansing
2. **Work with Missing data**
3. Detect and Handle Outliers

>>> Data analysis (to u, to technical, to)

>>> feature eng. ,, select the most important feature based of Domain Knowledge (u ,)

4. Deal with Imbalanced classes
5. Work with Categorical data
6. Feature Scaling
7. Split data to Train and Test Sets


First at all ... a closer look at your data



- What does the data model look like?
- What is the features distribution?
- What are the features with missing or inconsistent values?
- What are the most predictive features?
- Conduct a Exploratory Data Analysis (EDA)



Exploratory Data Analysis (EDA)




Exploratory Data Analysis (EDA)

is **a crucial step** in the machine learning process that involves visually and statistically analyzing the data

to gain insights, understand its characteristics, and identify patterns or trends,

And to make informed decisions about data preprocessing, feature engineering, and model selection

Exploratory Data Analysis overview



Why EDA is important in the machine learning process?

In general, EDA is a critical step in the machine learning process as it helps in :

- Understanding the data
- Assessing data quality
- Guiding feature engineering
- Informing model selection
- Effectively communicating results

EDA use case on clothing reviews

<https://www.kaggle.com/code/wonduk/eda-lstm-classification-on-clothing-reviews>

More EDA use case

<https://www.kaggle.com/wonduk/code>

Data Preprocessing

Feature Transformations

- ⬡ Data Cleaning or Cleansing
- ⬡ Work with Missing data
- ⬡ Detect and Handle Outliers
- ⬡ Deal with Imbalanced classes
- ⬡ Work with Categorical data
- ⬡ Feature Scaling
- ⬡ Split data to Train and Test Sets

Feature Engineering and Extraction

- Domain knowledge features (age , wh, **BMI**)
- Date and Time features
- String operations
- Web Data
- Geospatial features

Feature Transformations

1. **Data Cleaning or Cleansing**
2. Work with Missing data
3. Detect and Handle Outliers
4. Deal with Imbalanced classes
5. Work with Categorical data
6. Feature Scaling
7. Split data to Train and Test Sets


Data Cleaning or Cleansing

Cleaning data is the process of preparing the dataset for analysis. It is very important because the accuracy of machine learning or data mining models are affected because of poor quality of data.

So, data scientists spend a large amount of their time cleaning the dataset and transform them into a format with which they can work with. In fact, data scientists spend 80% of their time cleaning the data.

Example of problems

- ⬡ Columns can have missing data indicators like xx and ?.
- ⬡ Height column can have value of 0.
- ⬡ Weight column can have negative values.
- ⬡ ...



```
head  
dtypes
```

```
info >> nulls  
for num >> describe  
for cat >> value_count
```

Data Constraints Problems:

- **Data Type Constraints:**

- **Problem:** Incorrect data types in columns.
- **Solutions:** Convert to the correct data type.
- **Example in Action:** A revenue_usd column that contains values like "\$1,000" as strings instead of numeric data types.

- **Data Range Constraints:**

- **Problem:** Values outside the expected range.
- **Solutions:** Check for typos, drop rows, set to min/max, or impute.
- **Example in Action:** A gpa column with values exceeding the valid range, such as a GPA of 5.2.

- **Uniqueness Constraints:**

- **Problem:** Duplicates in the data.
- **Solutions:** Keep one, merge non-exact duplicates, or infer categories.
- **Example in Action:** Duplicate rows with the same name and phone_number columns but different height_cm values.

Text and Categorical Data Problems:

- **Membership Constraints for Categorical Data:**
 - **Problem:** Inconsistent categories.
 - **Solutions:** Drop rows, remap categories, or infer correct categories.
 - **Example in Action:** Two different entries for "New York" "NYC" , "New york" in the city column, causing inconsistency.
- **Length Violation for Text Data:**
 - **Problem:** Text length doesn't match expectations.
 - **Solutions:** Drop rows or set affected observations to missing.
 - **Example in Action:** A US phone_number column with entries that have 9 characters instead of the expected 14.
- **Text Data with Inconsistent Formatting:**
 - **Problem:** Text data not following a standard format.
 - **Solutions:** Drop rows or set affected observations to missing.
 - **Example in Action:** Inconsistent date formats like "dd-mm-yyyy" and "mm-dd-yyyy" in a date_of_birth column.

Data Cleaning (First Step in Preprocessing)

"Garbage in → Garbage out"

Why Cleaning First?

- Raw data is always messy.
- ML models are **sensitive** to errors.
- Cleaning builds the **foundation** for all later steps.
- Example: You can't cook with dirty vegetables → first wash & clean.

Data Problems Detection – Quick Checklist

1. **Inspect structure** → `shape`, `info`, `head()`
2. **Find missing values** → `isnull().sum()` (we see later)
3. **Detect duplicates** → `duplicated().sum()`
4. **Spot outliers** → `describe()`, `boxplot()` (we see later)
5. **Check category consistency** → `unique()`
6. **Catch invalid values** → logical rules (`Age < 0`, `Salary < 0`)
7. **Clean text fields** → `str.strip()`, `str.lower()`



Always **diagnose your dataset** before training

→ "Bad data in → bad model out."

→ "Garbage in → Garbage out"

Quick Data Inspection

- **df.shape** **# rows & columns**
 - Tells you (rows, columns).
 - Example: (2000, 10) → 2000 records, 10 features.
 - Helps to know if dataset is small, medium, or big.
- **df.head()** **# first rows**
 - Shows sample records.
 - Detect strange values early.
 - Example: Last row has Age = 999 → suspicious.
- **df.info()** **# data types**
 - Shows column names, data types (int, float, object, datetime).
 - Detects mismatches:
 - Age stored as "object" instead of "int".
 - Date stored as string instead of datetime.
- **df.describe()** **# summary stats**
 - Mean, min, max, std for each numeric column.
 - Quickly spot impossible values (e.g., Height min = -10).

Like a **doctor's first check-up**: weight, temperature, blood pressure.

Before deep analysis, you need to know the "overall health" of the dataset.



Why Important?

- Like a **doctor's first check-up**: weight, temperature, blood pressure.
- Before deep analysis, you need to know the "overall health" of the dataset.

Feature Transformations

1. Data Cleaning or Cleansing
2. **Work with Missing data**
3. Detect and Handle Outliers

>>> Data analysis (to u, to technical, to)

>>> feature eng. ,, select the most important feature based of Domain Knowledge (u ,)

4. Deal with Imbalanced classes
5. Work with Categorical data
6. Feature Scaling
7. Split data to Train and Test Sets

Work with Missing data

```
1 df.isnull().sum()
```



- ⬡ **Check Missing Data**
- ⬡ Drop Missing Data
- ⬡ Fill Missing data with Pandas
- ⬡ Fill Missing data with Sklearn SimpleImputer
- ⬡ Fill Missing data with Sklearn KNNImputer

```
[99]: df_with_nan= pd.DataFrame({'A':[1,2, None], 'B':[None, None, 6]})
df_with_nan
```

```
[99]:
```

	A	B
0	1.0	NaN
1	2.0	NaN
2	NaN	6.0

```
[101]: df_with_nan.isnull()
```

```
[101]:
```

	A	B
0	False	True
1	False	True
2	True	False

```
[109]: df_with_nan.isna()
```

```
[109]:
```

	A	B
0	False	True
1	False	True
2	True	False

```
[103]: df_with_nan.isnull().sum()
```

```
[103]: A    1
      B    2
      dtype: int64
```

```
•[105]: #total sum of null
df_with_nan.isnull().sum().sum()
```

```
[105]: 3
```

Check missing data

In Pandas missing data is represented by two value: None/NaN

we use a function isnull() to check even if NAN is true/false

we use a function isnull() to check number of nan in each column

```
[3] import pandas as pd
#read csv file
data=pd.read_csv('/content/data-2010-12-01.csv')
data.head(6)
```

	X	Y	Z	W
0	10	12.0	1.0	5.0
1	27	NaN	5.0	NaN
2	20	39.0	0.0	2.0
3	17	11.0	NaN	NaN
4	93	15.0	6.0	28.0
5	59	0.0	NaN	NaN

```
data.isnull()
```

	X	Y	Z	W
0	False	False	False	False
1	False	True	False	True
2	False	False	False	False
3	False	False	True	True
4	False	False	False	False
5	False	False	True	True

```
data.isnull().sum()
```

X	0
Y	1
Z	2
W	3
dtype:	int64

Drop Missing Data

1. When the data goes missing on 70–80 percent of the variable:

dropping the variable should be considered

2. When the data goes missing on 1–5 percent of the variable

dropping missing values only should be considered

`dropna( axis=0, how="any", subset=None, inplace=False)`

axis: possible values are {0 or 'index', 1 or 'columns'}, default 0. If 0, drop rows with null values. If 1, drop columns with missing values.

how: possible values are {'any', 'all'}, default 'any'. If 'any', drop the row/column if any of the values is null. If 'all', drop the row/column if all the values are missing.

subset: specifies the rows/columns to look for null values.

inplace: a boolean value. If True, the source DataFrame is changed and None is returned

Dropping rows/columns with at least 1 null value

```
[3] import pandas as pd
#read csv file
data=pd.read_csv('/content/data-2010-12-01.csv')
data.head(6)
```

	X	Y	Z	W
0	10	12.0	1.0	5.0
1	27	NaN	5.0	NaN
2	20	39.0	0.0	2.0
3	17	11.0	NaN	NaN
4	93	15.0	6.0	28.0
5	59	0.0	NaN	NaN

`data.dropna(axis=0)`

	X	Y	Z	W
0	10	12.0	1.0	5.0
2	20	39.0	10.0	2.0
4	93	15.0	6.0	28.0

Rows

`data.dropna(axis=1)`

	X
0	10
1	27
2	20
3	17
4	93
5	59

columns

Work with Missing data

```
1 df.dropna(axis = 0)
2
3 df.dropna(axis = 1)
```



- ⬡ Check Missing Data
- ⬡ **Drop Missing Data**
- ⬡ Fill Missing data with Pandas
- ⬡ Fill Missing data with Sklearn SimpleImputer
- ⬡ Fill Missing data with Sklearn KNNImputer

Fill Missing data with Pandas

1. Filling the missing data with the **mean or median** value if it's a **numerical variable**.
2. Filling the missing data with the **median** value only if it's a **numerical variable with outliers**
3. Filling the missing data with **mode** if it's a **categorical value**.

Work with Missing data

```
1 # fill with mean
2 df['BuildingArea'].fillna(df['BuildingArea'].mean(), inplace=True)
3
4
5 # fill with median
6 df['YearBuilt'].fillna(df['YearBuilt'].median(), inplace=True)
7
8
9 # fill with mode (most frequent)
10 df['Car'].fillna(df['Car'].mode()[0], inplace=True)
```

- ⬡ Check Missing Data
- ⬡ Drop Missing Data
- ⬡ **Fill Missing data with Pandas**
- ⬡ Fill Missing data with Sklearn SimpleImputer
- ⬡ Fill Missing data with Sklearn KNNImpute
- ⬡ r

Work with Missing data

```
1 from sklearn.impute import SimpleImputer
2
3 # define the Imputer properties
4 imputer = SimpleImputer(strategy='mean')    # can use 'median' or 'most_frequent'
5
6 # impute
7 df['BuildingArea'] = imputer.fit_transform(df[['BuildingArea']])
8
9 # get the filled value
10 imputer.statistics_
```

- ⬡ Check Missing Data
- ⬡ Drop Missing Data
- ⬡ Fill Missing data with Pandas
- ⬡ **Fill Missing data with Sklearn SimpleImputer**
- ⬡ Fill Missing data with Sklearn KNNImputer

عدد الغرف، المساحة ، السعر

???

120

3

Work with Missing data

لما الداتا تكون عندي صغيرة ؟

ليش ؟

$O(n^2)$

لازم الفيتشرز تكون مترابطة مع

بعض


```
import pandas as pd
from sklearn.impute import KNNImputer
df = pd.read_csv('melb_data.csv')

imputer = KNNImputer()
Num_df = df.select_dtypes(include = "number")
Cat_df = df.select_dtypes(include = "object_")
Num_df = pd.DataFrame(imputer.fit_transform(Num_df) , columns= Num_df.columns)

df = pd.concat([Num_df , Cat_df] , axis = 1)
```



- Check Missing Data
- Drop Missing Data
- Fill Missing data with Pandas
- Fill Missing data with Sklearn SimpleImputer
- Fill Missing data with Sklearn KNNImputer



Feature Transformations


- Data Cleaning or Cleansing
- Work with Missing data
- **Detect and Handle Outliers**
- Deal with Imbalanced classes
- Work with Categorical data
- Feature Scaling
- Split data to Train and Test Sets

Detect and Handle Outliers

In statistics, an outlier is a data point that differs significantly from other observations.


An outlier may be due to variability in the measurement or it may indicate experimental error, the latter are sometimes excluded from the data set.

An outlier can cause serious problems in statistical analyses.




Detect Outliers

```
1 sns.boxplot(x='DIS', data=df)  
2 sns.stripplot(x='DIS', data=df)
```



- ⬡ **Detect Outliers with Visualization.**
- ⬡ Detect and Handle Outliers with IQR.
- ⬡ Handle Outliers with business value



Detect and Handle Outliers

```
1 from datascist.structdata import detect_outliers
2
3 # remove outliers
4 outliers_indices = detect_outliers(df, 0, df.columns)
5 df.drop(outliers_indices, inplace=True)
6
7
8 # replace outliers with median value for each column
9 for col in df.columns:
10     outliers_indices = detect_outliers(df, 0, [col])
11     col_median = df[col].median()
12     df[col].iloc[outliers_indices] = col_median
```

- Detect Outliers with Visualization.
- **Detect and Handle Outliers with IQR.**
- Handle Outliers with business value

Detect and Handle Outliers with Business value

- Essentially, instead of removing outliers from the data.
- you change their values to something more representative of your data set , value related to the business of your data , that called business value.
- For example the “**Alice in Wonderland**” movie duration in this dataset is 557 Minutes which is an outlier , so we can replace this outlier value with its real duration 75 Minute
- Detect Outliers with Visualization.
- Detect and Handle Outliers with IQR.
- Handle Outliers with business value**

	title	year	certificate	runtime_min	genre
490	Alice in Wonderland	1951	PG-13	557	Animation, Adventure, Comedy
262	The Beatles: Get Back	2021	PG-13	468	Documentary, Biography, Music
167	Guardians of the Galaxy Vol. 2	2017	PG-13	23	Action, Adventure, Comedy
474	Willow	2022-	PG-13	46	Action, Adventure, Drama
14	Eternals	2021	PG-13	156	Action, Adventure, Fantasy
...	...	...	...	...	...
628	All About Steve	2009	PG-13	99	Comedy, Romance
638	Ghost Rider: Spirit of Vengeance	2011	PG-13	96	Action, Fantasy, Thriller
706	Kim Possible	2019 TV Movie	PG-13	86	Action, Adventure, Comedy
1355	In the Mix	2005	PG-13	95	Comedy, Crime, Drama
348	The Wolverine	2013	PG-13	0	Action, Sci-Fi

#	Method	When to Use	When NOT to Use
1	Detect (IQR / Z-score)	First step in EDA to find extreme values	Not a solution by itself
2	Ignore	Outliers are real and rare with low impact	When they strongly affect mean or model
3	Remove	Clear data errors (typos, sensor faults)	When values are valid business cases
4	Cap / Clip	Keep all rows, reduce extreme influence	When outliers are key business targets
5	Transformation	Skewed data (sales, revenue, finance)	Many zeros or negative values
6	Separate Analysis	Outliers represent a special segment	When no business meaning exists

Feature Transformations

- Data Cleaning or Cleansing
- Work with Missing data
- Detect and Handle Outliers

>>> Data analysis (to u, to technical, to)



>>> feature eng. ,, select the most important feature based of Domain Knowledge (u ,)

- **Deal with Imbalanced classes**
- Work with Categorical data
- Feature Scaling
- Split data to Train and Test Sets

Deal with Imbalanced classes

Imbalanced classes are a common problem in machine learning classification where there are an unbalanced ratio of observations in each class. Class imbalance can be found in many different areas including medical diagnosis, spam filtering, and fraud detection.

3200 rows



Deal with Imbalanced classes

Data Leakage



```
1 from imblearn.under_sampling import RandomUnderSampler
2
3 x = df.drop('Class', axis=1)
4 y = df['Class']
5
6 rus = RandomUnderSampler()
7
8 x_train, y_train = rus.fit_sample(x_train, y_train)
```

Downsampling


A strategy to handle imbalanced classes by creating a random subset of the majority of equal size to the minority class.

ChrisAlbon

- ⬡ Down-sampling or Under-sampling Majority Class
- ⬡ Generate Synthetic Samples with SMOTE

Deal with Imbalanced classes

```
1 from imblearn.over_sampling import SMOTE
2
3 x = df.drop('Class', axis=1)
4 y = df['Class']
5
6 smote = SMOTE()
7
8 x_train, y_train = smote.fit_sample(x_train, y_train)
```



- ⬡ Down-sampling or Under-sampling Majority Class
- ⬡ Generate Synthetic Samples with SMOTE

Deal with Imbalanced classes (Combination)

- Oversampling methods duplicate or create new synthetic examples in the minority class, whereas under sampling methods delete or merge examples in the majority class.
- Both types of resampling can be effective when used in isolation, although can be more effective when both types of methods are used together.

Deal with Imbalanced classes



```
# Manually Combine Over and Undersampling Methods
from imblearn.over_sampling import RandomOverSampler from
imblearn.under_sampling import RandomUnderSampler from imblearn.pipeline import
Pipeline
#over = RandomOverSampler(sampling_strategy=0.1) over =
SMOTE(sampling_strategy=0.1) under = RandomUnderSampler(sampling_strategy=0.5)
pipeline = Pipeline(steps= [('o', over), ('u', under)]) # fit and apply the
pipeline X_resampled, y_resampled = pipeline.fit_resample(X, y)
```

Feature Transformations

- Data Cleaning or Cleansing
- Work with Missing data
- Detect and Handle Outliers
- Deal with Imbalanced classes
- **Feature Scaling**
- Work with Categorical data
- Split data to Train and Test Sets

Feature Scaling

- It is common to scale data prior to fitting a machine learning model.
- This is because data often consists of many different input variables or features (columns) and each may have a different range of values or units of measure, such as feet, miles, kilograms, dollars, etc.
- If there are input variables that have very large values relative to the other input variables, these large values can dominate or skew some machine learning algorithms.
- The result is that the algorithms pay most of their attention to the large values and ignore the variables with smaller values.
- This includes algorithms that use a weighted sum of inputs :

1. like linear regression

age >> 18 to 70

2. logistic regression

salary >> 10000 to 500000

3. artificial neural networks

as well as algorithms that use distance measures between examples such as :

1. k-nearest neighbors

2. support vector machines.

Feature Scaling

```
1 from sklearn.preprocessing import MinMaxScaler
2
3 scaler = MinMaxScaler()
4
5 scaler.fit(x_train)
6
7 scaled_x_train = scaler.transform(x_train)
8 scaled_x_test = scaler.transform(x_test)
```

MINMAX SCALING

Rescales feature values to
between 0 and 1

Original value x_i →

Minimum value in feature $\min(x)$ →

Maximum value in feature $\max(x)$ →

Rescaled value $x'_i = \frac{x_i - \min(x)}{\max(x) - \min(x)}$

ChrisAlbon

- ⬡ Normalization with Sklearn MinMaxScaler.
- ⬡ Standardizing data with StandardScaler.

Feature Scaling

```
1 from sklearn.preprocessing import StandardScaler
2
3 scaler = StandardScaler()
4
5 scaler.fit(x_train)
6
7 scaled_x_train = scaler.transform(x_train)
8 scaled_x_test = scaler.transform(x_test)
```

STANDARDIZATION

Standardized feature value $X'_i = \frac{\text{Value of the } i\text{th observation } X_i - \text{Mean of the feature vector } \bar{X}}{\text{Standard deviation of the feature vector } \sigma}$

Standardization is a common scaling method. X'_i represents the number of standard deviations each value is from the mean value. It rescales a feature to have a mean of 0 and unit variance.

Chris Albon

- Normalization with Sklearn MinMaxScaler.
- Standardizing data with StandardScaler.

Normalization with Sklearn MinMaxScaler.

- it is the simplest method and consists of rescaling the range of features to scale the range in $[0, 1]$

Standardizing data with StandardScaler.

- Feature standardization makes the values of each feature in the data have zero mean and unit variance.

Feature Scaling (Robust Scalar)

- Sometimes an input variable may have outlier values. These are values on the edge of the distribution that may have a low probability of occurrence yet are overrepresented for some reason.
- **Outliers can skew a probability distribution and make data scaling using standardization difficult as the calculated mean and standard deviation will be skewed by the presence of the outliers.**
- One approach to standardizing input variables in the presence of outliers is to ignore the outliers from the calculation of the mean and standard deviation, then use the calculated values to scale the variable. This is called robust standardization or **robust data scaling**.
- This can be achieved by calculating the median (50th percentile) and the 25th and 75th percentiles.
- The values of each variable then have their median subtracted and are divided by the interquartile range (IQR) which is the difference between the 75th and 25th percentiles

$$\text{value} = \frac{\text{value} - \text{median}}{p_{75} - p_{25}}$$

```
from sklearn.preprocessing import RobustScaler  
scaler = RobustScaler(quantile_range=(25.0, 75.0))  
X_train_rb = scaler.fit_transform(X_train)  
X_test_rb = scaler.transform(X_test)
```

✗ Common Mistakes (Feature Scaling):

1. **Applying Feature Scaling on the entire dataset *before* splitting into train and test sets (75-80% train , 20-25% test)**
→ This can cause **data leakage**, as information from the test set may influence the scaling parameters (like mean and standard deviation).
2. **Forgetting to apply the same scaler (fitted on training data) to the test data**

⚠ Potential Mistakes (Imbalanced Classes):

1. **Misusing Under-sampling → Loss of valuable information!**
→ Removing too many samples from the majority class can result in losing important patterns in the data.
2. **Using SMOTE before splitting the data → Data Leakage!**
→ If you apply SMOTE before splitting, synthetic samples are generated using information from the entire dataset, including the test set — which causes leakage.

Feature Transformations

- Data Cleaning or Cleansing
- Work with Missing data
- Detect and Handle Outliers
- Deal with Imbalanced classes
- Feature Scaling
- **Work with Categorical data**
- Split data to Train and Test Sets


Work with Categorical ordinal data (Label Encoder)

- Label encoding converts the data in machine-readable form, but it assigns a unique number(starting from 0) to each class of data , in a sequence ordinal manner.
- To use label encoding, your data should must have inherent order in its logic
- We can't use it with nominal data (non – ordered).
- An attribute having output classes Mexico, Paris, Dubai. On Label Encoding, this column lets Mexico is replaced with 0, Paris is replaced with 1, and Dubai is replaced with 2.
- With this, it can be interpreted that Dubai has high priority than Mexico and Paris while training the model, But actually, there is no such priority relation between these cities here.

```
1 size_dict = {'XS':1,  
2             'S':2,  
3             'M':3,  
4             'L':4,  
5             'XL':5,  
6             'XXL':6}  
7  
8 # apply using map  
9 df['Size'] = df['Size'].map(size_dict)
```

- **Work with Ordinal Features with pandas map method.**


Work with Categorical Nominal data (ONE-HOT Encoder / Binary Encoding)



```
1 df = pd.get_dummies(df, columns=['Color', 'Brand'], drop_first=True)
```

- Work with Ordinal Features with pandas map method.
- **Work with Nominal Features with pandas get_dummies method.**
- Work with Nominal Features with sklearn OneHotEncoder method.

Work with Categorical Nominal data (ONE-HOT Encoder / Binary Encoding)



- Work with Ordinal Features with pandas map method.
- Work with Nominal Features with pandas get_dummies method.
- **Work with Nominal Features with sklearn OneHotEncoder method.**

Binary Encoding for Nominal Data

- The categories are encoded in Ordinal Numeric form.
- Then those integers are converted into binary code, so for example 5 becomes 101 and 10 becomes 1010
- Then the digits from that binary string are split into separate columns.
- Each observation is encoded across the columns in its binary form.
- Example: if we have Country Feature with 16 different countries:
 - By One-Hot Encoding it will result in 16 columns.
 - By Binary Encoding, it will result only in 4 columns

```
X = pd.DataFrame(list('abcdefghijklmnop'))
X
```

	0
0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

```
from sklearn.preprocessing import OneHotEncoder
ohe = OneHotEncoder(sparse=False)
countries= ohe.fit_transform(X.values.reshape(-1,1)).astype(int)
pd.DataFrame(countries)
```

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
1	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0
2	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0
3	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0
5	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0
6	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0
7	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0
8	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0
9	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0	0
10	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0	0
11	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0	0
12	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0	0
13	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0	0
14	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	0
15	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1

```
import category_encoders as ce

ce_bin = ce.BinaryEncoder()
ce_bin.fit_transform(X)
```

	0_0	0_1	0_2	0_3	0_4
0	0	0	0	0	1
1	0	0	0	1	0
2	0	0	0	1	1
3	0	0	1	0	0
4	0	0	1	0	1
5	0	0	1	1	0
6	0	0	1	1	1
7	0	1	0	0	0
8	0	1	0	0	1
9	0	1	0	1	0
10	0	1	0	1	1
11	0	1	1	0	0
12	0	1	1	0	1
13	0	1	1	1	0
14	0	1	1	1	1
15	1	0	0	0	0

45 class > 6 col
 $= 2^? = 45$
 $= 2^6 = 64$

31 class > one hot > 31 clo
 31 class > binary enc > 5 clo

1 2 4 8 16
 1 1 1 1 > binary number for 15
 1 0 0 1 > binary number for 9
 0 0 1 0 1 > binary number for 20

Feature Transformations

- Data Cleaning or Cleansing
- Work with Missing data
- Detect and Handle Outliers
- >>>> feature eng
- >>> split data if u need feature scaling or Imbalanced data
- Deal with Imbalanced classes
- Feature Scaling
- Work with Categorical data
- **Split data to Train and Test Sets**

Split data to Train and Test Sets

When you're working on a model and want to train it, you obviously have a dataset. But after training, we have to test the model on some test dataset.

the obvious solution is to split the dataset you have into two sets, one for training and the other for testing; and you do this before you start training your model.

	weight	height	drinks alcohol	healthy	
0	112	181	0	0	
1	123	165	1	1	
2	176	167	1	1	
3	145	154	1	1	← Y_train
4	198	181	0	0	
5	211	202	1	0	
6	145	201	1	1	
7	181	153	1	1	
8	90	142	0	1	← Y_test
9	101	169	1	1	

```
1 from sklearn.model_selection import train_test_split
2
3 x = df.drop('healthy', axis=1)
4 y = df['healthy']
5
6 x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2)
```

first 90% form data with label “Yes”

Next 10% form data with label “No”

split 80/20

<<< 80% هيتدرب على yes فقط؟؟ انو مشافش no وهو بيتدرب
<< 20% نصهم No فبالتالي المودل مش هيقدر يكتشفهم ليش؟ لانو
ما شافهم اثناء التدريب

>> shuffle = True

train > 90% yes, 10% no

test > 90% yes, 10% no

stratify = y

Sampling techniques in Train and Test Sets

Important in imbalanced data.

Parameter of train_test_split:

1. Shuffle

With shuffle=True, data is shuffled then the split happens.

2. Stratify

Makes a split so that a proportion of the values in the sample produced will be the same as the proportion of values provided to parameter stratify.

3. Random state (random_state = 42) (0,1,7,42,1234)

If added you ensure that the split remains the same whenever you run the code

	weight	height	drinks alcohol	healthy
0	112	181	0	0
1	123	165	1	1
2	176	167	1	1
3	145	154	1	1
4	198	181	0	0
5	211	202	1	0
6	145	201	1	1
7	181	153	1	1
8	90	142	0	1
9	101	169	1	1

[illegible]